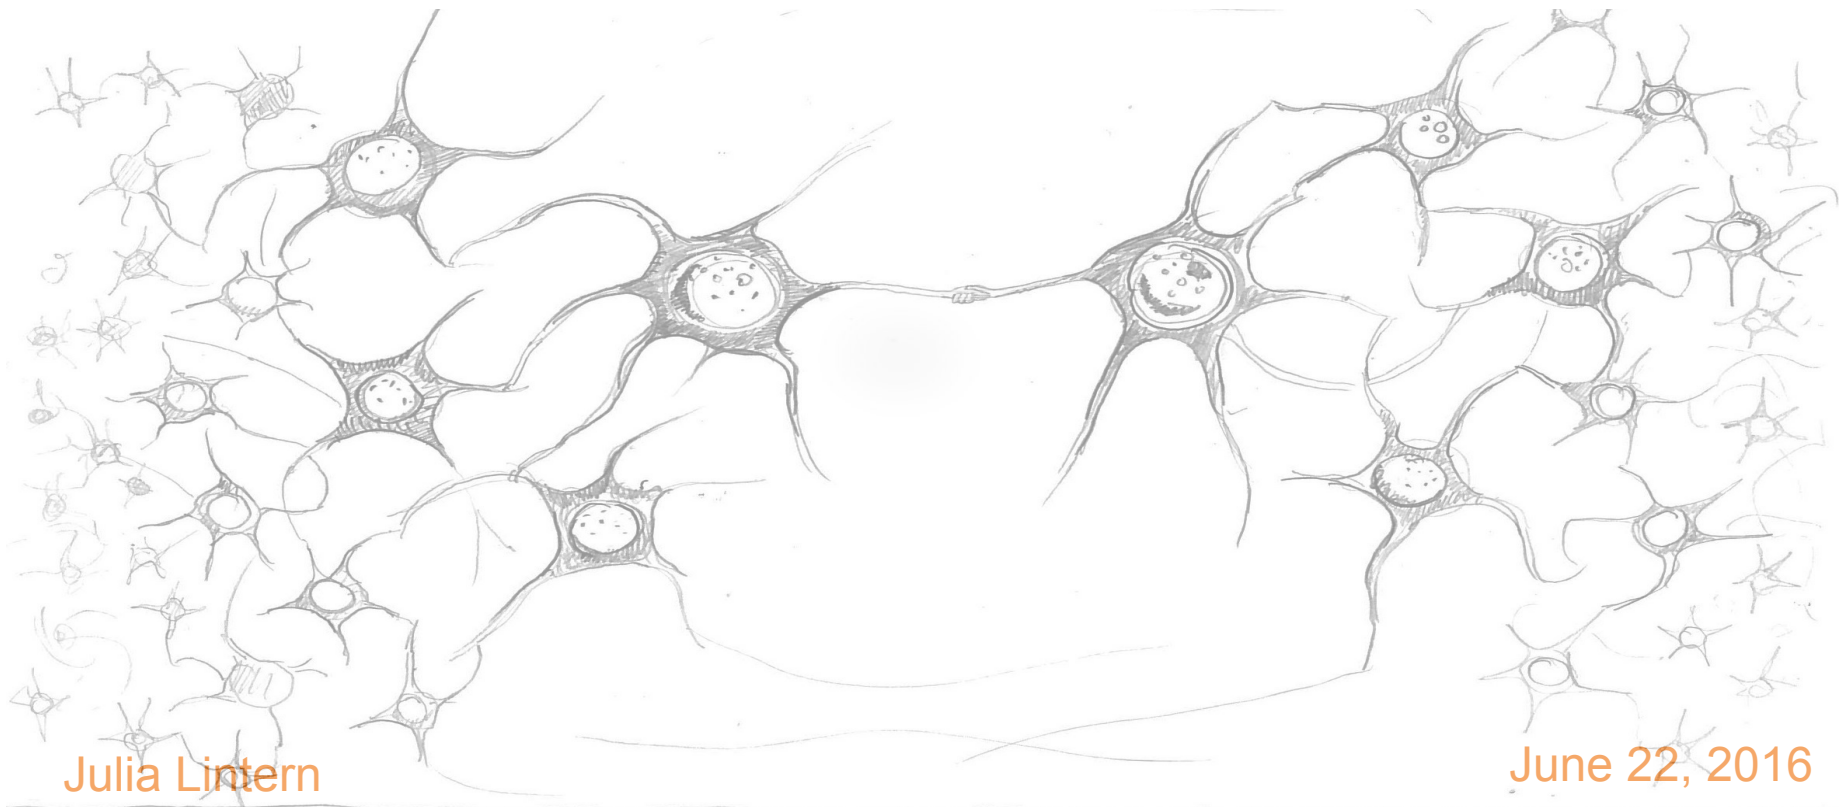


# Navigating Neural Nets

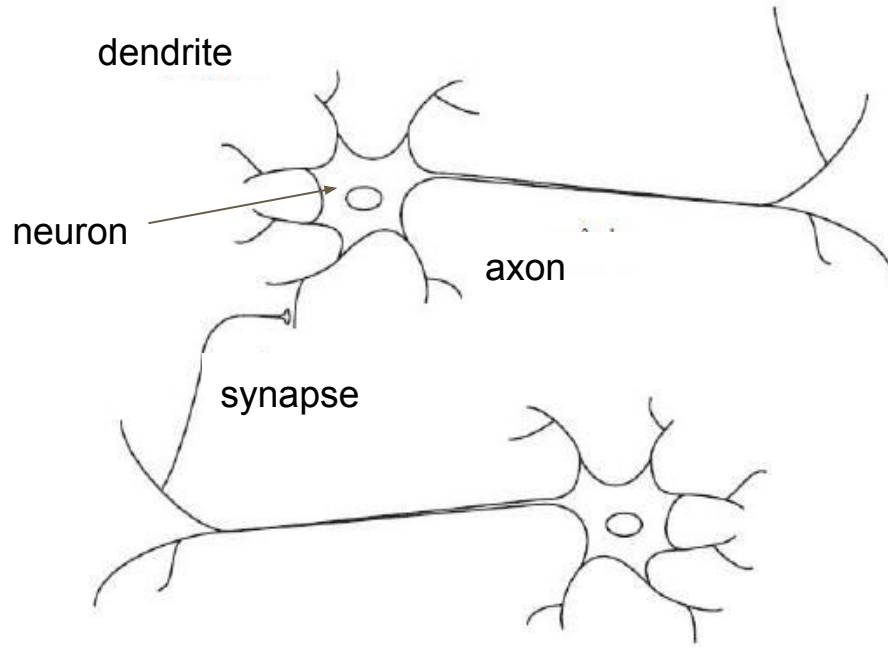


Julia Lintern

June 22, 2016

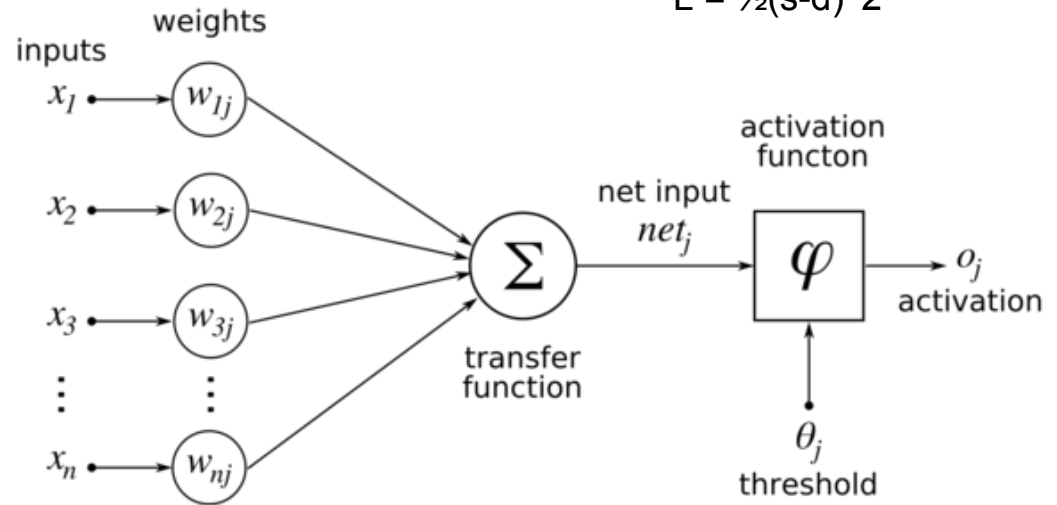
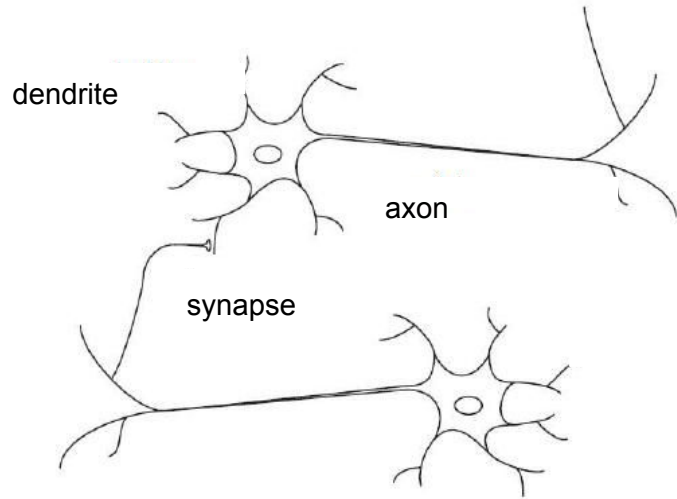
# The Brain Analogy

(our cartoon neuron)



# The Brain Analogy

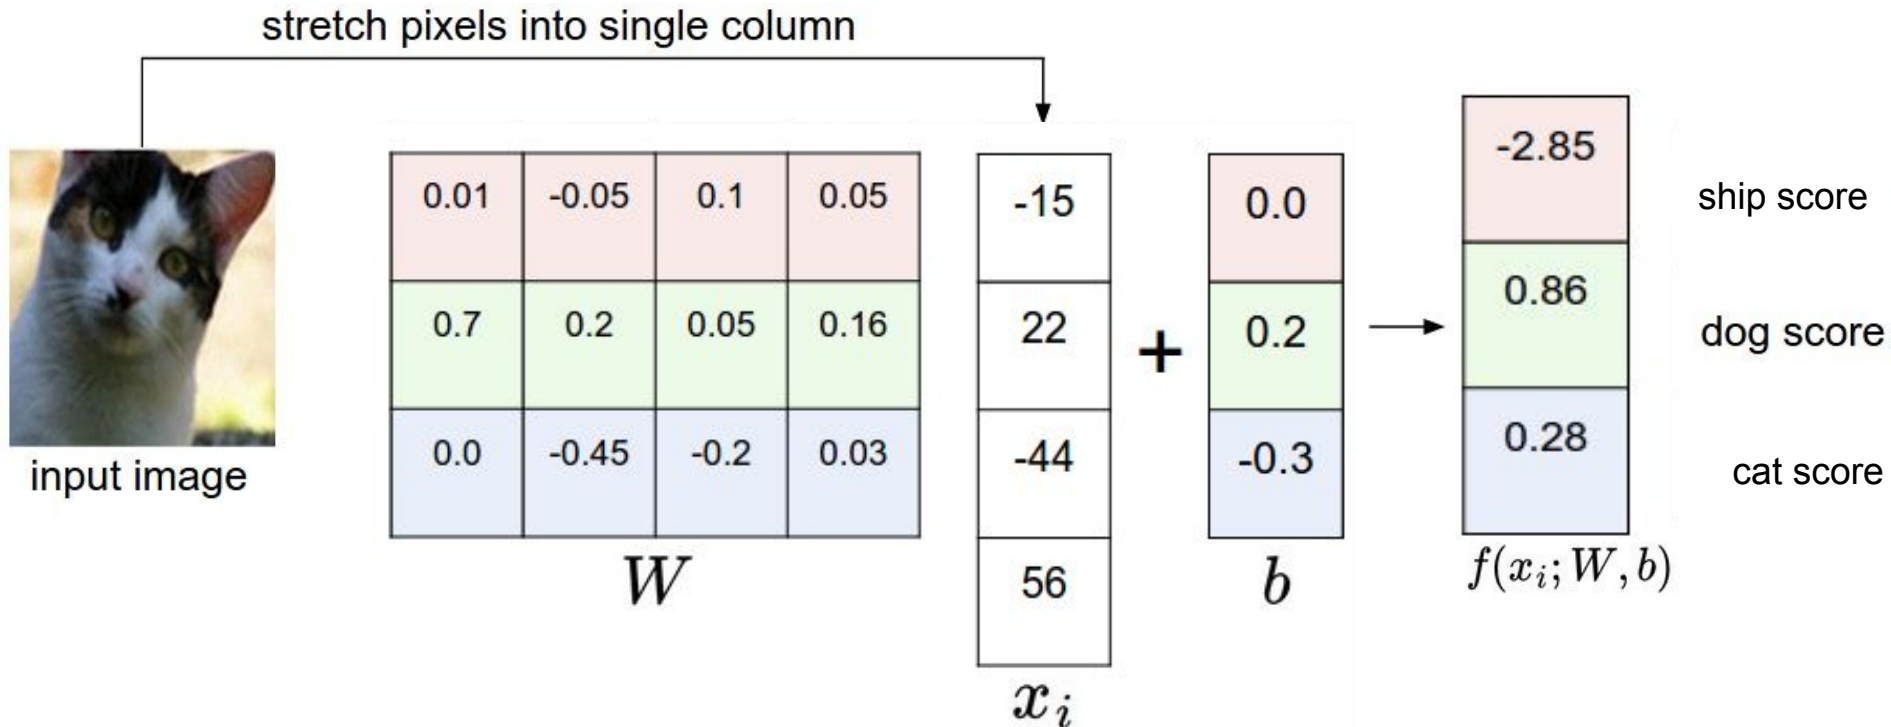
(cartoon neuron & mathematical neuron)



$$s = f(x, w)$$

$$L = \frac{1}{2}(s - d)^2$$

# The Linear Classifier Analogy



## Losses: Multiclass SVM (Hinge) Loss

0.01	-0.05	0.1	0.05
0.7	0.2	0.05	0.16
0.0	-0.45	-0.2	0.03

$W$

-15
22
-44
56

$x_i$

+

0.0
0.2
-0.3

$b$

-2.85
0.86
0.28

ship score

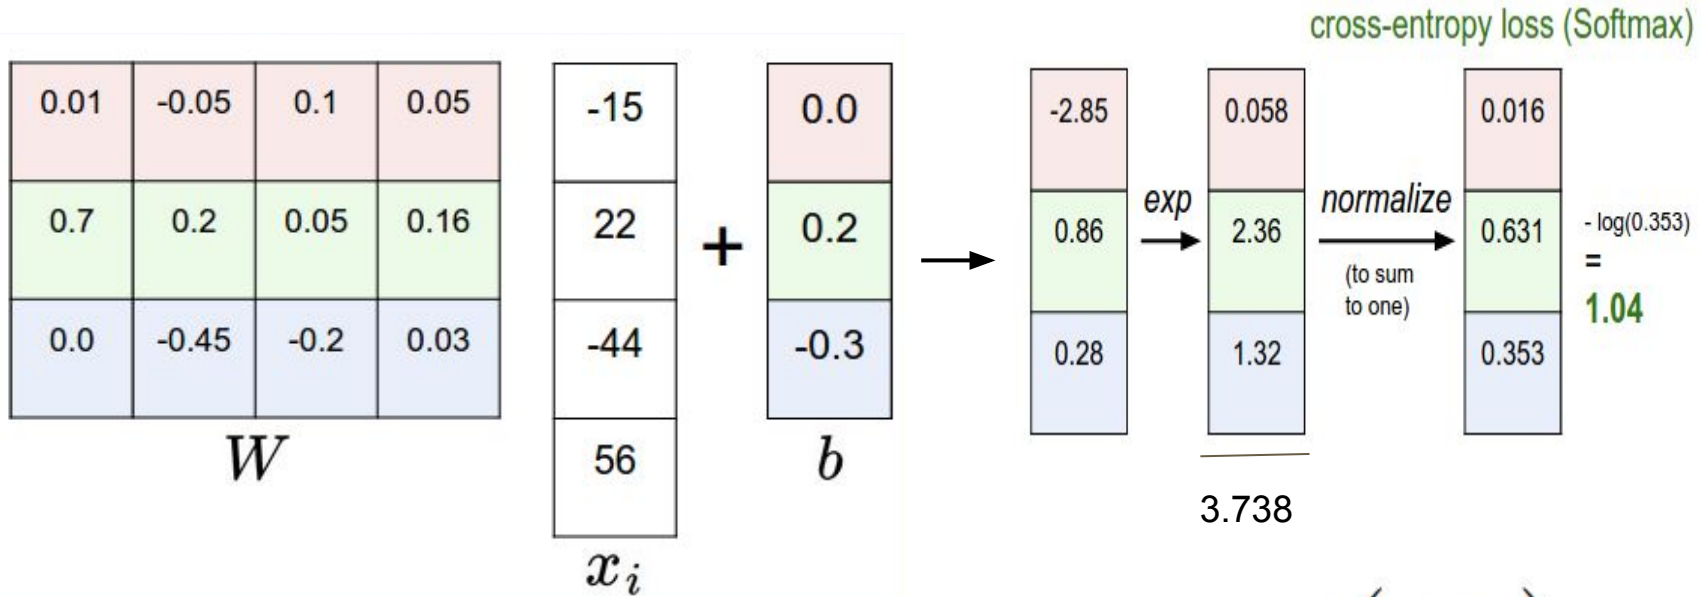
dog score

cat score

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + \Delta)$$

$$= \max(0, -2.85 - 0.28 + \Delta) \stackrel{\Delta=1}{=} 0 + \max(0, 0.86 - 0.28 + \Delta) \stackrel{\Delta=1}{=} 0 + 1.58$$

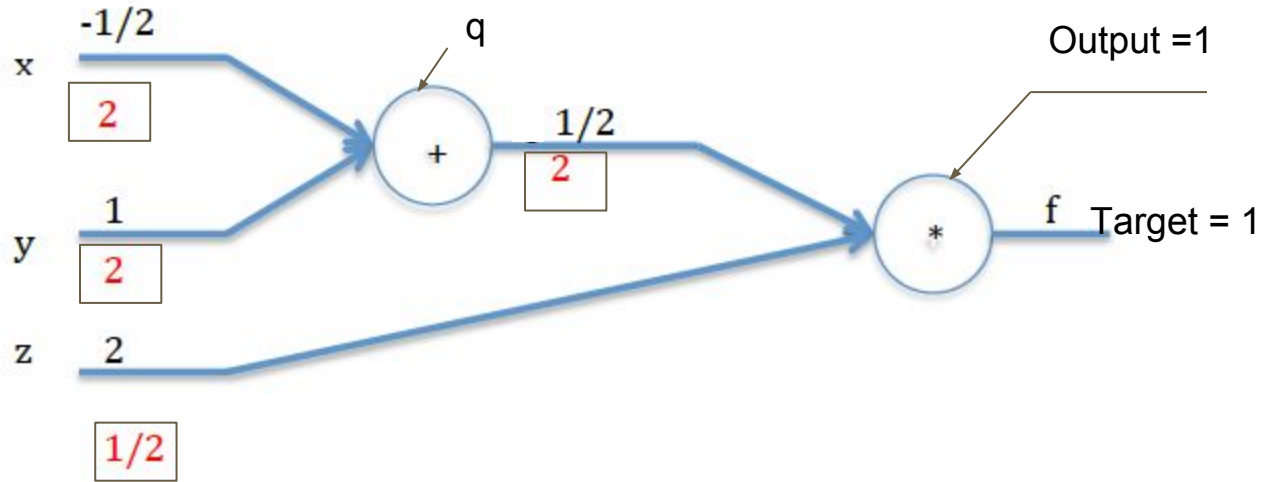
## Losses: Softmax (Cross-Entropy) Loss



Softmax:  $f_j(z) = \frac{e^{z_j}}{\sum_k e^{z_k}}$

Cross-Entropy  $L_i = -\log\left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}}\right)$

# BackPropagation



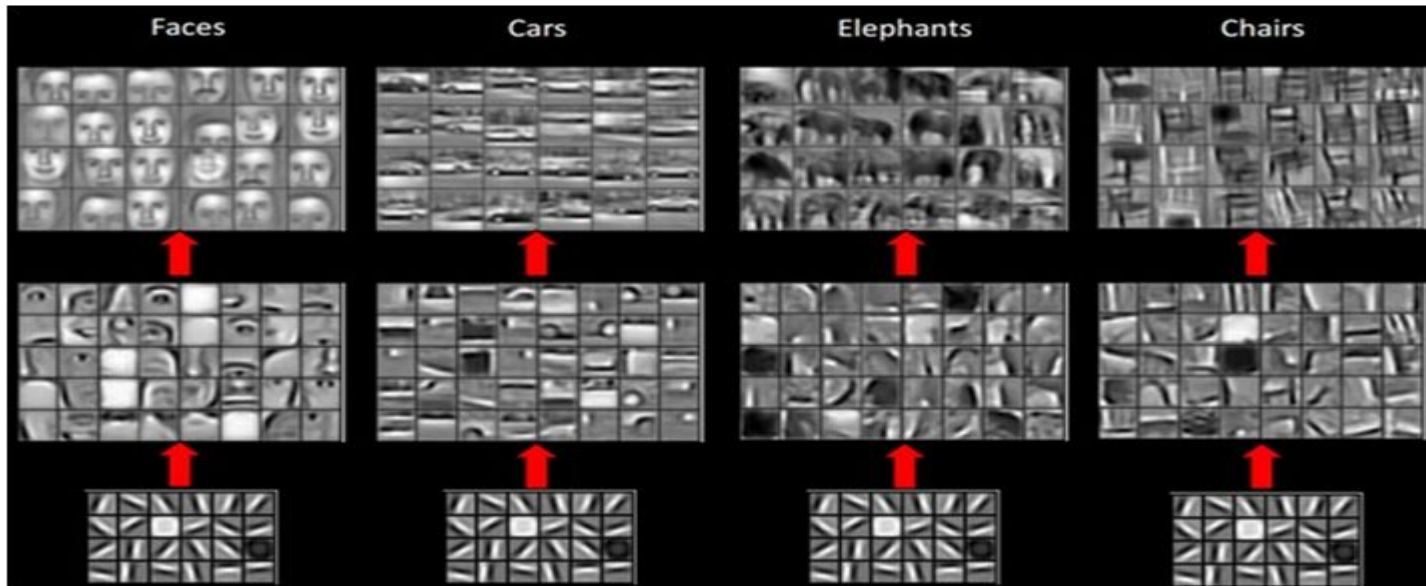
$$f(x,y,z) = (x+y)z$$
$$q = (x+y)$$

$$\frac{df}{dx} = ?$$
$$\frac{df}{dy} = ?$$

$$\Rightarrow \text{chain rule: } \frac{df}{dx} = \frac{df}{dq} \left( \frac{dq}{dx} \right)$$
$$\Rightarrow \text{chain rule: } \frac{df}{dy} = \frac{df}{dq} \left( \frac{dq}{dy} \right)$$

# Convolutional Neural Nets

Very similar to Neural Nets.. But how are they different ?

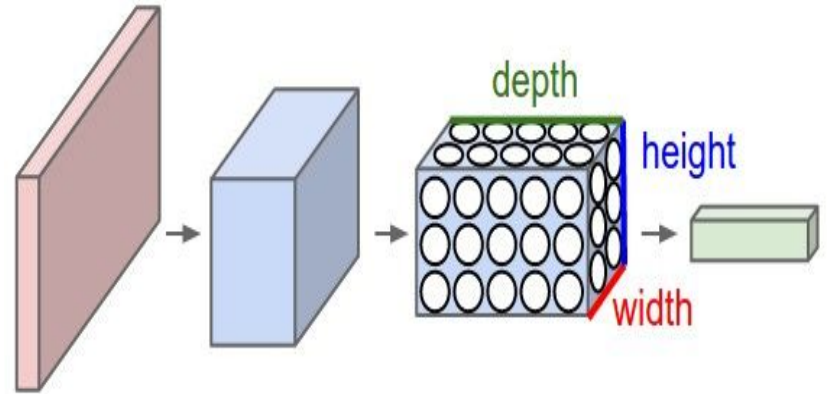
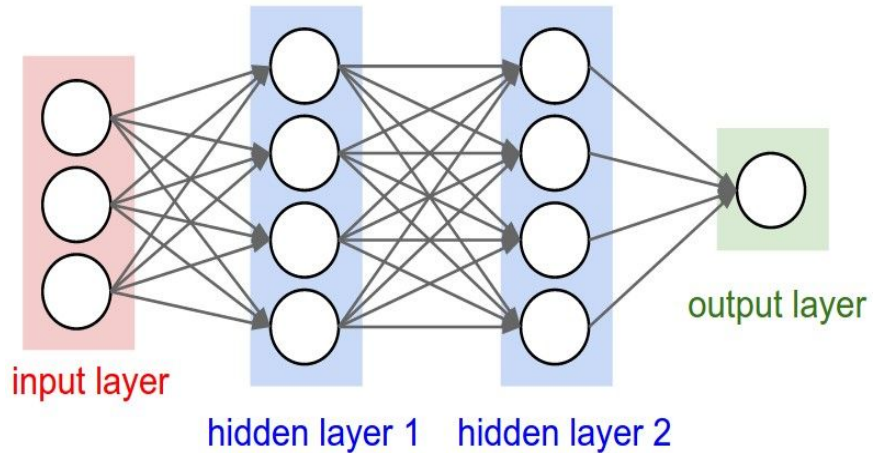




# Convolutional Neural Nets

## Vs. Neural Nets

- Input is an image: Leverage 3D Structure
- Fully Connected ? Not entirely!



# The CNN Family

## Winners of the ILSVRC ImageNet challenges

**AlexNet (2012, Krizhevsky):** Popularized CNNs - 1st to incorporate consecutive convolutional layers

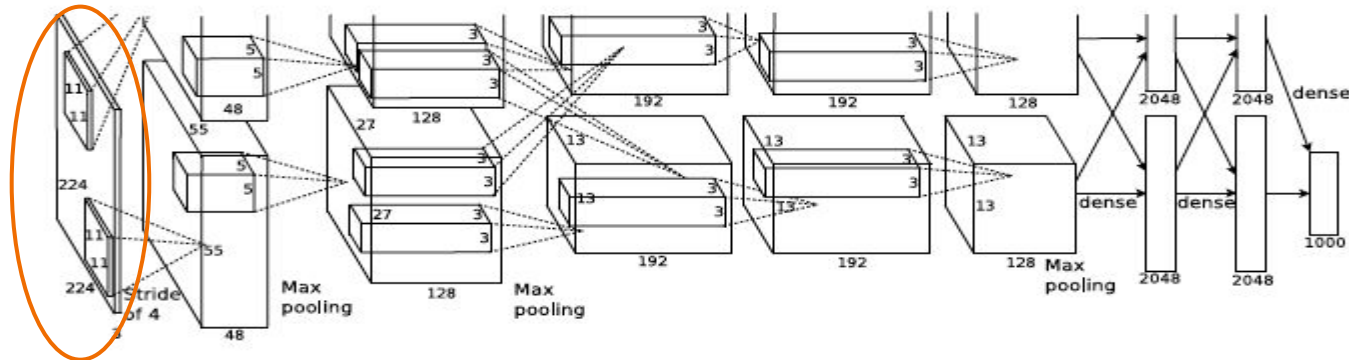
**GoogLeNet / Inception (2014, Szegedy):** Drastically reduced the # of parameters used (from 60 million to 4 million)

**ResNet (2015, Kaiming He):** Residual Network : famous for skip-connections and heavy use of batch-normalization; also removes some fully connected layers (at end of network)



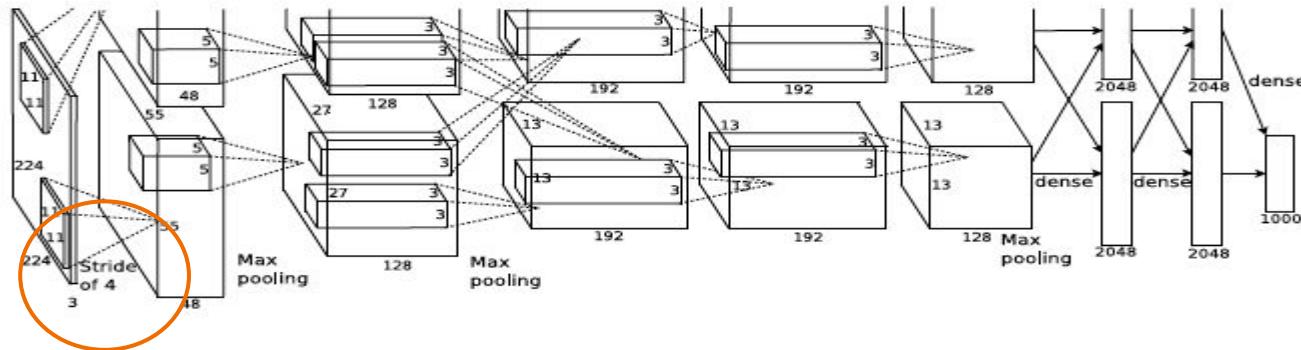
# Convolutional Neural Nets : Architecture

- 1) **Input Layer:** Raw pixel values of the image  
(ex:  $224 \times 224 \times 3$  (3 ~ color channels (RGB)))
- 2) Conv Layer
- 3) Pool Layer
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)



# Convolutional Neural Nets : Architecture

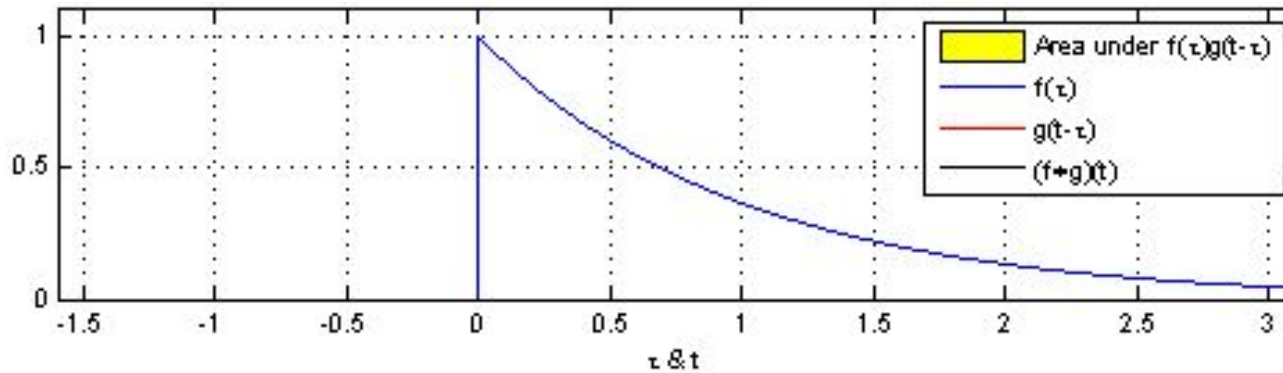
- 1) Input Layer: Raw pixel values of the image
- 2) **Conv Layer: Dot product between weights and the small region of input volume (ex: 11 x 11 x 3 filters)**
- 3) Pool Layer
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)



# Convolutional Neural Nets : Architecture

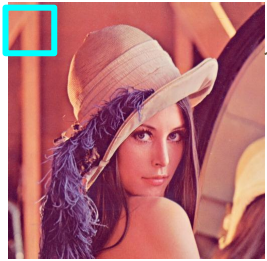
## What is a Convolution?

$$f * g = \int f(t - \tau)g(\tau)d\tau$$



# Convolutional Neural Nets : Architecture

## What is a Convolution?



1 1 1 0  
1 1 1 0  
1 1 1 0

1 1 1  
1 -8 1  
1 1 1

→ 0

1 1 1 0  
1 1 1 0  
1 1 1 0

1 1 1  
1 -8 1  
1 1 1

→ -3

# Convolutional Neural Nets : Architecture

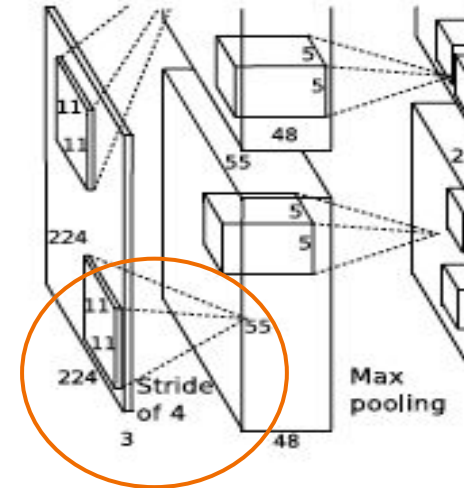
## What is a Convolution?

**Convolutional Layer:**  $(W-F + 2P)/S + 1$

- **W** : Input Volume size
- **F**: Receptive Field size of the Conv Layer Neuron
- **P**: Zero- Padding
- **S**: Stride

$$(224 - 11 + 2(3))/4 + 1 = 55$$

Conv Layer Output ~ 55 x 55 x 96 (ie : 55<sup>2</sup> neurons in each layer)





# Convolutional Neural Nets : Architecture

What is a Convolution?

Voila. We have 96 filters.





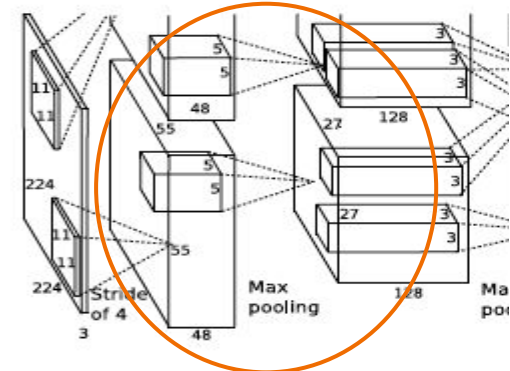
# Convolutional Neural Nets : Architecture

- 1) Input Layer
- 2) Conv Layer
- 3) **Pooling Layer: Performs downsampling operation**
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)

Our Eqn :  $O = (W - F) / S + 1$

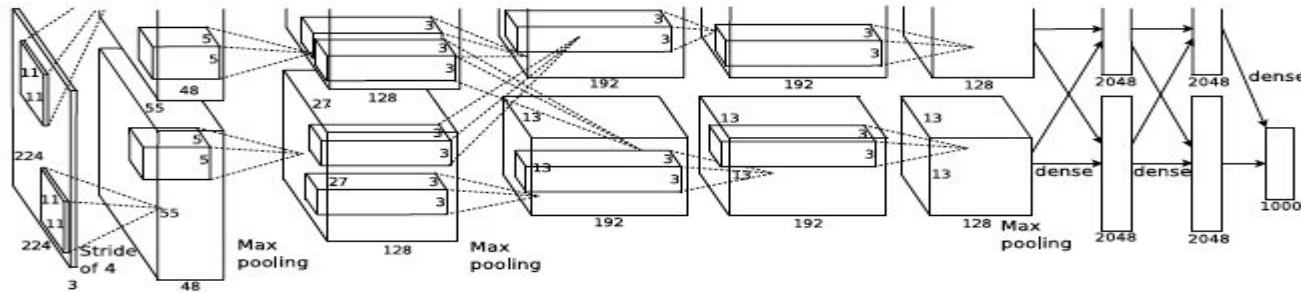
AlexNet: use 3 x 3 MaxPooling w/ stride = 2

$$O = (55 - 3) / 2 + 1 = 27$$



# Convolutional Neural Nets : Architecture

- 1) Input Layer: Raw pixel values of the image
- 2) Conv Layer:
- 3) Pool Layer:
- 4) **ReLU Layer: Apply an elementwise activation function**  
(ex :  $\max(0, x)$  thresholding output dimension ~ same as input)
- 5) FC (Fully Connected Layer)



\*The ReLU non-linearity is applied to the output of every convolutional and fully-connected layer.

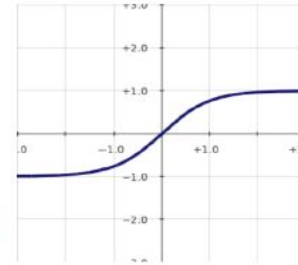
# Convolutional Neural Nets : Architecture

## ReLU Layer:

Traditionally:

$f(x) = \tanh(x)$  or  $fx = (1+e^{-x})^{-1}$  ( Very slow to train)

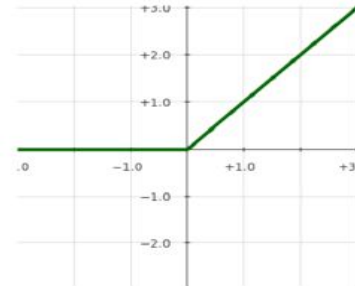
$f(x) = \tanh(x)$



Now:

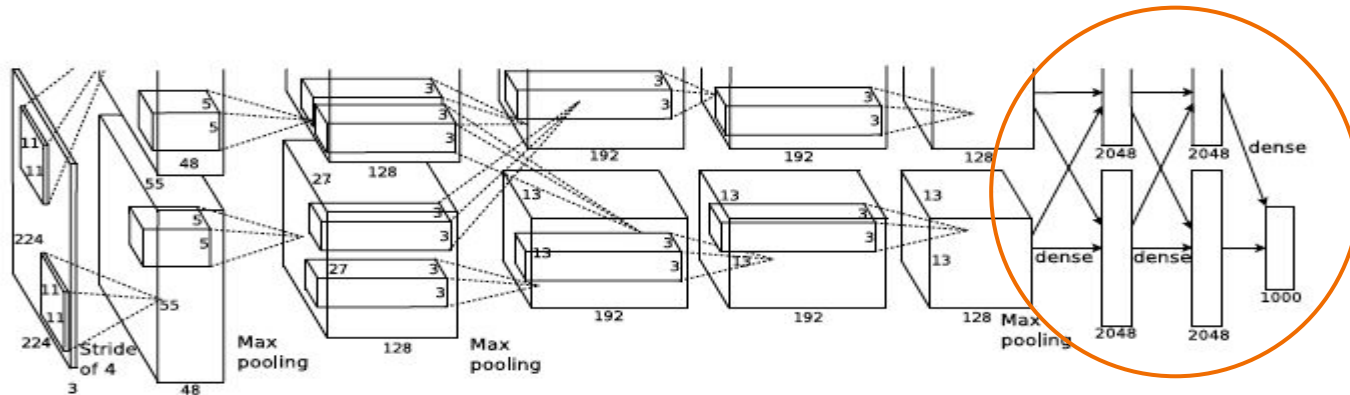
$f(x) = \max(0, x)$  (Faster to train )

$f(x) = \max(0, x)$



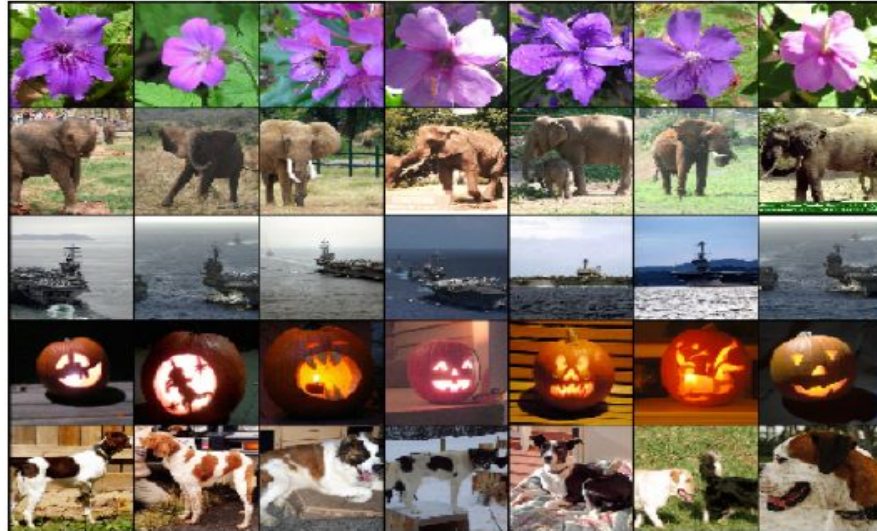
# Convolutional Neural Nets : Architecture

- 1) Input Layer: Raw pixel values of the image
- 2) Conv Layer:
- 3) Pool Layer:
- 4) ReLU Layer:
- 5) **FC (Fully Connected) Layer** : Each neuron will be connected to all activations of the previous volume. The output layer will compute class scores (ex:  $[1 \times 1 \times 1000]$  )



# Convolutional Neural Nets : Architecture

Output from the final 4096 fully connected layer :



## *Torch:*

- Fast. Easy to integrate with GPUs
- Many modular pieces that are easy to combine <https://github.com/imagenet-multiGPU.torch/blob/master/models/alexnet.lua>
  - Written in Lua



## *TensorFlow:*

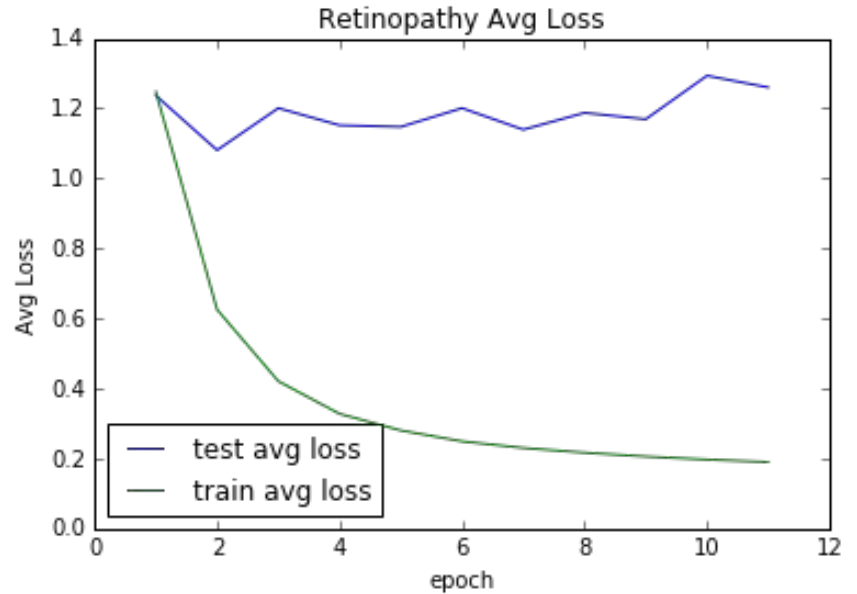
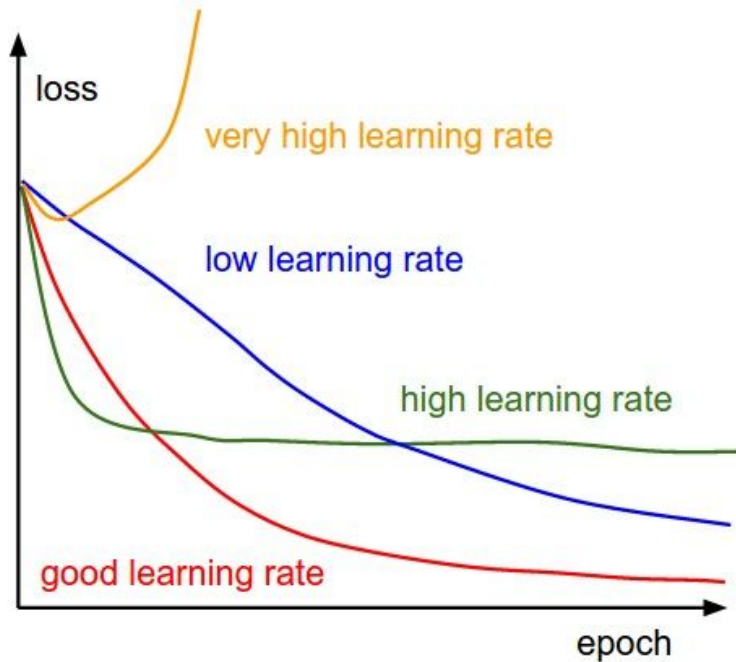
- Written in python & numpy
- Tensorboard for visualization
  - Latest releases can be buggy
  - Can be many x slower than Torch



<https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#LaunchInstanceWizard:>

# Learning from the Learning Process

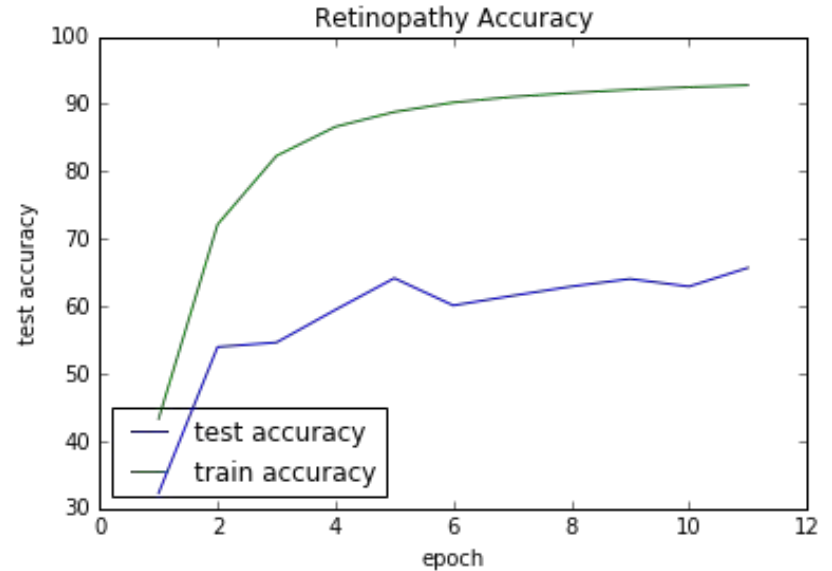
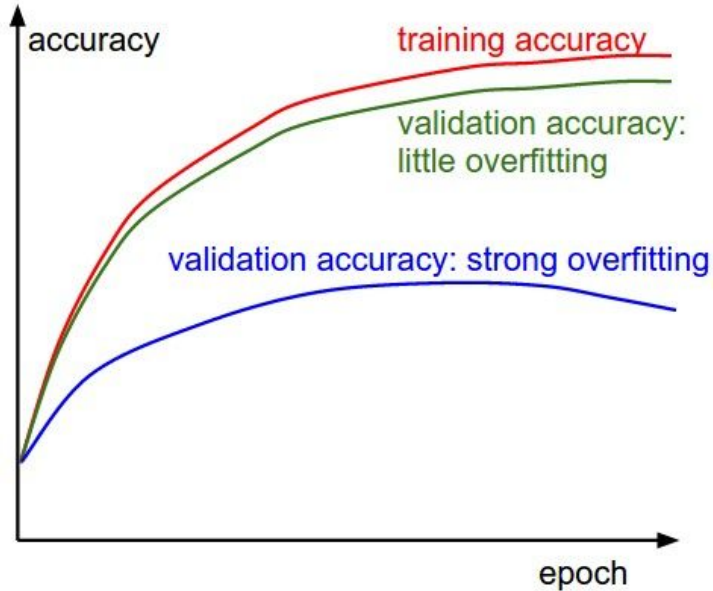
## 1) Loss functions



\* Tip: Change the learning rate!

# Learning from the Learning Process

## 2) Accuracy



\* Tip: Increase L2 weight penalty , Increase Drop-Out, More Data (possibly with jitter) - -try batch norm ?



# Learning from the Learning Process

## 3) First-layer Visualizations

Visualized weights from the 1st layer of the network:  
(smooth, diverse features indicate that training is going well)

